

1 Introduction

2 Image Preprocessing

2.1 Convolution

For two continuous functions f and g , the convolution is defined as

$$(f * g)(x) = \int_{-\infty}^{\infty} g(\tau)f(x - \tau)d\tau \quad \text{being commutative, associative and linear.}$$

2.2 Fourier Transform

The Fourier transform of f is given by:

$$\hat{f}(k) := \int_{-\infty}^{\infty} f(x)e^{-2\pi i k x} dx \quad \text{With the inverse Fourier transform:}$$

$$f(x) := \int_{-\infty}^{\infty} \hat{f}(k)e^{2\pi i k x} dk$$

2.3 Impulse Functions

$$\text{Dirac impulse: } \delta(x) = \begin{cases} 0 & \text{if } x \neq 0 \\ \infty & \text{if } x = 0 \end{cases} \quad \text{with } \int_{-\infty}^{\infty} \delta(x)dx = 1$$

$$\text{And a series of Dirac impulses: } \sum_{n=-\infty}^{\infty} \delta(x - nT)$$

2.4 Nyquist Sampling Theorem

If f is band bounded with cutoff frequency k_0 $\hat{f}(k) = 0$ for all k with $|k| \geq k_0$
then it is completely determined by giving its ordinates at a series of points spaced at most $\frac{1}{2k_0}$
meaning the sample frequency must be larger than $2k_0$

2.5 Quantisation

Characteristic with equidistant steps of size Δ

$$g(x) = \max\{0, \min\{g_{\max}, \left\lfloor \frac{I(x)}{\Delta} + \frac{1}{2} \right\rfloor\}\}$$

$$h(x) = \Delta g(x) - \frac{\Delta}{2}$$

yielding an error of non-overdriven quantiser:

$$I(x) - h(x) \in \left[-\frac{\Delta}{2}, \frac{\Delta}{2}\right]$$

2.6 Gamma Correction

$$g_{out} = g_{max} \left(\frac{g_{in}}{g_{max}} \right)^\gamma$$
 which keeps black and white, is a nonlinear transformation

2.7 Models of Blur

Blur can be modeled with convolution

$$g_{blurred} = g_{sharp} * p$$
 with p modeling the blur

$$p_{motion}(x) = \begin{cases} \frac{1}{n} & \text{if } -n < x \leq 0 \\ 0 & \text{otherwise} \end{cases}$$
 along x-axis by n pixels

$$p_{Gauss}(x) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{1}{2} \frac{x^2}{\sigma^2}}$$

2.8 Wiener Deconvolution

Restore a sharp image from a blurred image with a Wiener filter: $g_{blurred} = g_{sharp} * p + v$ with p being a point-spread function (blur), v being pixel noise and assuming g_{sharp} and v being independent.

With $g_{restored} = f * g_{blurred}$, find optimal f that minimizes

$$e(k) = E \left[|\hat{g}_{sharp}(k) - \hat{g}_{restored}(k)|^2 \right]$$
 with E being the expectation value.

2.9 Models of Noise

Statistical noise: $g_{noisy}(x, y) = g_{sharp}(x, y) + v(x, y)$

Malfunctioning sensors: $g_{noisy}(x, y) = \begin{cases} g_{sharp}(x, y) & \text{with probability } p \\ \text{arbitrary} & \text{otherwise} \end{cases}$

3 Edge-Detection

3.1 Finding Edges

Approximating a derivative by difference: $\frac{\partial g}{\partial x} \approx \frac{g(x+1) - g(x-1)}{2}$

3.2 Laplace Operator

The 2D analogon to the second order derivatie is the Laplace operator:

$$\nabla^2 g = \frac{\partial^2 g}{(\partial x)^2} + \frac{\partial^2 g}{(\partial y)^2} \text{ which can be appriximated with}$$

$$\nabla^2 g \approx g(x+1, y) + g(x-1, y) + g(x, y+1) + g(x, y-1) - 4g(x, y)$$

Because the second order derivative is very noisy, so combine Laplacian with Gaussian smoothing:

$$\nabla^2 (G * g) = (\nabla^2 G) * g \text{ which yields}$$

$$\nabla^2 G = \frac{x^2 + y^2 - 2\sigma^2}{\sigma^4} G(x, y) \text{ called Laplacian of Gaussian (LoG) or Mexican Hat}$$

The Laplacian of Gaussian can be approximated by calculating the Difference of Gaussian (DoG)

$$DoG(x, y) = G_{\sigma_1}(x, y) - G_{\sigma_2}(x, y)$$

3.3 Corner Detection

Using dissimilarity measure to find patches of maximal dissimilarity for local moves:

$$d := \sum_{(u,v) \in rectangle} (g(u + \Delta u, v + \Delta v) - g(u, v))^2 \approx \begin{pmatrix} \Delta u \\ \Delta v \end{pmatrix}^T \begin{pmatrix} \sum (\frac{\partial g}{\partial u})^2 & \sum \frac{\partial g}{\partial u} \frac{\partial g}{\partial v} \\ \sum \frac{\partial g}{\partial u} \frac{\partial g}{\partial v} & \sum (\frac{\partial g}{\partial v})^2 \end{pmatrix} \begin{pmatrix} \Delta u \\ \Delta v \end{pmatrix}$$

For special choice of coordinate system it becomes a diagonal matrix:

$$d := \begin{pmatrix} \Delta u \\ \Delta v \end{pmatrix}^T \begin{pmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{pmatrix} \begin{pmatrix} \Delta u \\ \Delta v \end{pmatrix} = \lambda_1 (\Delta u)^2 + \lambda_2 (\Delta v)^2 \text{ w.l.o.g. } \lambda_1 \geq \lambda_2 \geq 0$$

4 Curve Fitting

4.1 2D Geometry of Lines

Line in normal form: $0 = \langle \vec{n}, \vec{x} \rangle + c$ so the distance of a point \vec{r} from the line:

$$d = \|\vec{l} - \vec{r}\| = |\langle \vec{n}, \vec{r} \rangle + c|$$

4.2 Hough Transform

Every line in 2D can be represented as: $x \cos \phi + y \sin \phi + c = 0$ so there is a 2D space of all line represented by (ϕ, c) called Hough Space.

4.3 Total Least Squares

To estimate a line through a set of points we need to search the line parameters that minimise d_i

$$\text{minimise } \sum_{i=1}^N d_i^2 \text{ subject to } \langle \vec{n}, \vec{n} \rangle = 1 \quad \text{Using the Lagrange function:}$$

$$L(\vec{n}, c, \lambda) = \sum_{i=1}^n d_i^2 - \lambda(\langle \vec{n}, \vec{n} \rangle - 1) \quad \text{then zeroing the partial derivative with respect to}$$

c and then to n , we can arrive at

$$\begin{pmatrix} \alpha & \beta \\ \beta & \gamma \end{pmatrix} \vec{n} = \lambda \vec{n} \quad \text{hence } \lambda \text{ is Eigenvalue and } \vec{n} \text{ is Eigenvector. We get two solutions}$$

of Eigenvalue Problem: $\lambda_1 \geq \lambda_2 \geq 0$ where λ_2 minimises distances and λ_1 maximises distances

4.4 Weighted Least Squares

Because of outliers introduce weights $w_i \geq 0$, one for each edge point, then solve:

$$\text{minimise } \sum_{i=1}^N w_i d_i^2 \text{ subject to } \langle \vec{n}, \vec{n} \rangle = 1. \quad \text{But how to choose } w_i? \text{ Replace } d_i \text{ by a term}$$

that grows more slowly!

4.5 M-Estimators

M-Estimator means $\boxed{\text{minimise } \sum_{i=1}^N \rho(d_i) \text{ subject to } \langle \vec{n}, \vec{n} \rangle = 1}$ with a suitable function ρ

The derivatives of the Lagrange function of M-Estimators approach and Weighted Least Squares approach are equal if $\boxed{w_i = \frac{\frac{\partial \rho(d_i)}{\partial d_i}}{2d_i}}$. So running Weighted Least Squares with appropriate weights implements M-Estimators. E.g.:

$$\boxed{\rho_{Cauchy} : d \mapsto c^2 \log(1 + \frac{d^2}{c^2})}$$

$$\boxed{\rho_{Huber} : d \mapsto \begin{cases} d^2 & \text{if } |d| \leq k \\ 2k|d| - k^2 & \text{otherwise} \end{cases}}$$

4.6 Least Trimmed Sum of Squares Estimator (LTS)

$\boxed{\text{minimise } \sum_{i=1}^p d_{i:N} \text{ subject to } \langle \vec{n}, \vec{n} \rangle = 1}$ with $p < N$ and $d_{i:N}$ the i-th element in the ordered list of point-distances. p is typically given as a percentage of N .

4.7 RANSAC

Search a line that passes nearby as many points as possible.

$$\boxed{\text{minimise } \sum_{i=1}^N \sigma(d_i) \text{ with } \sigma(d_i) = \begin{cases} 0 & \text{if } |d_i| \leq \theta \\ 1 & \text{if } |d_i| > \theta \end{cases}}$$

4.8 Estimating Circles

Parametric representation of circles:

$$\boxed{(x - m_1)^2 + (y - m_2)^2 - r^2 = 0}$$

Euclidean distance of point (x,y) from the circle:

$$\boxed{d_E = \left| \sqrt{(x - m_1)^2 + (y - m_2)^2} - r \right|}$$

Algebraic distance:

$\boxed{d_A = |(x - m_1)^2 + (y - m_2)^2 - r^2|}$ Because minimising the Euclidean distance cannot be solved analytically, minimize the sum of algebraic distances. Use a weighted approach, chose w_i so that we implement Euclidean by using weighted Algebraic.

5 Color

6 Segmentation

6.1 Level Set Evolution

Two class segmentation can be represented by indicator function:

$$\phi(\vec{x}) \begin{cases} < 0 & \text{if pixel } \vec{x} \text{ belongs to segment} \\ > 0 & \text{if pixel } \vec{x} \text{ belongs to background} \end{cases} \text{ with } |\phi(\vec{x})| = \text{distance of } \vec{x} \text{ from contour}$$

Modeling temporal evolution of signed distance function with $\phi(\vec{x}, t)$. When tracking a point $\vec{x}(t)$ on the boundary over time, it obviously follows $\phi(\vec{x}(t), t) = 0 \text{ for all } t$

$$0 = \frac{d\phi(\vec{x}(t), t)}{dt} = \nabla \phi \frac{\partial \vec{x}}{\partial t} + \frac{\partial \phi}{\partial t}. \text{ Solving this for } \frac{\partial \phi}{\partial t} \text{ yields:}$$

$$\frac{\partial \phi}{\partial t} = - \nabla \phi \frac{\partial \vec{x}}{\partial t}$$

6.2 Mumford-Shah

Shrink contour if $(I - C_{in})^2 > (I - C_{out})^2$

Expand contour if $(I - C_{in})^2 < (I - C_{out})^2$

Mumford-Shah based segmentation:

$$\frac{\partial \vec{x}}{\partial t} = \frac{\nabla \phi}{\|\nabla \phi\|} (C + P - \lambda_1(I - C_{in})^2 + \lambda_2(I - C_{out})^2)$$

7 Camera Optics

7.1 World To Image Mapping

Point (x, y, z) is projected onto (x', y') via intercept theorem:

$$z \begin{pmatrix} x' \\ y' \end{pmatrix} = \begin{pmatrix} f & 0 \\ 0 & f \end{pmatrix} \begin{pmatrix} x \\ y \end{pmatrix}$$

Mapping from camera to image frame:

$$\begin{pmatrix} u \\ v \end{pmatrix} = \begin{pmatrix} \alpha & -\beta \cot \theta \\ 0 & \frac{\beta}{\sin \theta} \end{pmatrix} \begin{pmatrix} x' \\ y' \end{pmatrix} + \begin{pmatrix} u_0 \\ v_0 \end{pmatrix}$$

Mapping from world to camera frame:

$$\begin{pmatrix} x \\ y \\ z \end{pmatrix} = R \begin{pmatrix} \xi \\ \eta \\ \zeta \end{pmatrix} + \vec{t}$$

Rewriting all those, given (ξ, η, ζ) we can calculate (u, v) by

$$\begin{pmatrix} \tilde{x} \\ \tilde{y} \\ \tilde{z} \end{pmatrix} = A(R | \vec{t}) \begin{pmatrix} \xi \\ \eta \\ \zeta \\ 1 \end{pmatrix} \quad \text{and} \quad \begin{pmatrix} u \\ v \end{pmatrix} = \frac{1}{\tilde{z}} \begin{pmatrix} \tilde{x} \\ \tilde{y} \end{pmatrix}$$

The other way round, we can obtain (ξ, η, ζ) given (u, v) by

$$\begin{pmatrix} \xi \\ \eta \\ \zeta \end{pmatrix} = z R^T A^{-1} \begin{pmatrix} u \\ v \\ 1 \end{pmatrix} - R^T \vec{t} \quad \text{with } z \leq 0$$

7.2 Snell's law

$$n_e \sin \theta_e = n_t \sin \theta_t \quad \text{with} \quad n_{\text{medium}} = \frac{v_{\text{vacuum}}}{v_{\text{medium}}}$$

7.3 Thin lenses

Via Intercept theorem, we get the lens equation:

$$\frac{1}{f} = \frac{1}{b} + \frac{1}{g}$$

7.4 Depth of Field

Via intercept theorem with:

$$\boxed{\nabla g = g_{far} - g_{near} = 2 \frac{g d_h (g - f)}{d_h^2 - (g - f)^2}} \text{ with } \boxed{d_h = \frac{Df}{\epsilon}} \text{ as hyperfocal distance.}$$

7.5 Radial Distortion

Mathematical modeling with even polynomials:

$$\boxed{\begin{pmatrix} x_d \\ y_d \end{pmatrix} = (1 + k_1 r^2 + k_2 r^4) \begin{pmatrix} x' \\ y' \end{pmatrix}} \text{ with } r^2 = (x')^2 + (y')^2$$

7.6 Camera Calibration

Tsai's approach:

From world-to-image-mapping take $M := A(R | \vec{t})$ in which M is a 3x4 matrix.

Determine camera parameters by minimizing sum of squares, which means zeroing partial derivatives.

7.7 Telecentric Lenses

Magnification of telecentric lens: $\boxed{B = \frac{b - f}{f} G}$

Depth of field: $\boxed{\nabla g = 2 \frac{\epsilon}{D} (g - f)}$

8 Illumination

Irradiance of a surface point, meaning amount of incident light:

$$I(\theta_i, \phi_i)$$

Radiance of a surface point, meaning amount of emitted light:

$$I(\theta_e, \phi_e)$$

Total irradiance of a surface patch (in sperical coordinates):

$$I_0 = \int_0^{2\pi} \int_0^{\frac{\pi}{2}} I(\theta_i, \phi_i) \sin\theta_i \cos\theta_i d\theta_i d\phi_i$$

Reflectance of a surface patch can be modeled by the bidirectional reflectance distribution function (BRDF):

$f(\theta_e, \phi_e | \theta_i, \phi_i)$ describing how much light is radiated in direction (θ_e, ϕ_e) if one unit of light is irradiated from direction (θ_i, ϕ_i)

8.1 Reflectance

So total amount of radiated light:

$$I_0 = \int_0^{2\pi} \int_0^{\frac{\pi}{2}} f(\theta_e, \phi_e | \theta_i, \phi_i) I(\theta_i, \phi_i) \sin\theta_i \cos\theta_i d\theta_i d\phi_i$$

Lambertian Reflectance:

$$f(\theta_e, \phi_e | \theta_i, \phi_i) = \frac{1}{\pi} \text{ hence: } L(\theta_e, \phi_e) = \frac{1}{\pi} L_0$$

Specular Reflectance, mirroring effect:

$$f(\theta_e, \phi_e | \theta_i, \phi_i) = \frac{\delta(\theta_e - \theta_i) \delta(\phi_e - \phi_i - \pi)}{\sin\theta_i \cos\theta_i} \text{ hence: } L(\theta_e, \phi_e) = I(\theta_e, \phi_e - \pi) \text{ Combina-}$$

tions of Reflectance:

$$f(\theta_e, \phi_e | \theta_i, \phi_i) = \frac{\eta}{\pi} + (1 - \eta) \frac{\delta(\theta_e - \theta_i) \delta(\phi_e - \phi_i - \pi)}{\sin\theta_i \cos\theta_i}$$

9 3D Reconstruction

9.1 Stereovision by Triangulation

For camera 1:

$$\begin{pmatrix} \xi \\ \eta \\ \zeta \end{pmatrix} = z R^T A^{-1} \begin{pmatrix} u \\ v \\ 1 \end{pmatrix} - R^T \vec{t} \quad \text{with unknown depth } z$$

Same for camera 2 and equaling:

$z\vec{p} - \vec{q} = z'\vec{p}' - \vec{q}'$ is overdetermined, lines of sight might be skewed. Minimize square of norm instead!

9.2 Phase Shift

Measure gray value of pixel 4 times with phase shift multiples of $\frac{\pi}{2}$:

$$g_i = A \sin\left(\frac{u}{w} 2\pi + \dots\right) + B \quad \text{with wavelength } w. \text{ Solve for A and B.}$$

$$u = \frac{\text{atan2}(g_1 - g_3, g_2 - g_4)}{2\pi} w + kw \quad \text{only if } A \neq 0$$

10 Pattern Recognition

10.1 Bayesian Classification

assign new object to the class with the largest posteriori probability:

$$\hat{c} = \arg \max P(c \mid x_{new}) \quad \text{and} \quad P(c \mid x) \propto P(x \mid c)P(c)$$

10.2 Linear Classification

A linear classifier is a function that implements a function of the kind:

$$\vec{x} \mapsto \begin{cases} +1 & \text{if } \langle \vec{x}, \vec{w} \rangle + b \geq 0 \\ -1 & \text{otherwise} \end{cases} \quad \text{with } \vec{w} \text{ the weight vector and } b \text{ the bias weight.}$$

Margin, meaning the minimal distance between a hyperplane and the convex hull of the training patterns:

$$\rho = \min \left(d^{(i)} \frac{\langle \vec{x}^{(i)}, \vec{w} \rangle + b}{\|\vec{w}\|} \right)$$

10.3 Support Vector Machine

SVM is a linear classifier that maximizes the margin ρ . Training a SVM means solving:

$$\text{maximize } \rho^2 \quad \text{subject to} \quad d^{(i)} \frac{\langle \vec{x}^{(i)}, \vec{w} \rangle + b}{\|\vec{w}\|} \geq \rho \quad \text{for all } i$$

Where $\|\vec{w}\|$ is a degree of freedom, so set $\|\vec{w}\| = \frac{1}{\rho}$ for convenience. That leaves us to:

$$\text{minimize } \frac{1}{2} \|\vec{w}\|^2 \quad \text{subject to} \quad d^{(i)} (\langle \vec{x}^{(i)}, \vec{w} \rangle + b) \geq 1 \quad \text{for all } i$$

10.4 Fault-tolerant SVM

Soft margin SVM:

$$\text{minimize } \frac{1}{2} \|\vec{w}\|^2 + C \sum_i \xi_i \quad \text{subject to} \quad d^{(i)} (\langle \vec{x}^{(i)}, \vec{w} \rangle + b) \geq 1 - \xi_i \quad \text{for all } i$$

10.5 Nonlinear SVM

Assume nonlinear transformation:

$$\theta : \begin{cases} R^n \rightarrow R^m \\ \vec{x} \mapsto \theta(\vec{x}) = \vec{X} \end{cases}$$

Useful kernel-functions are:

$$K(\vec{x}, \vec{y}) = \langle \vec{x}, \vec{y} \rangle \quad \text{dot product}$$

$$K(\vec{x}, \vec{y}) = (\langle \vec{x}, \vec{y} \rangle)^d \quad \text{or} \quad (\langle \vec{x}, \vec{y} \rangle + 1)^d \quad \text{polynomial kernels}$$

$$K(\vec{x}, \vec{y}) = e^{-\frac{\|\vec{x} - \vec{y}\|^2}{2\sigma^2}}$$

10.6 Validation Process

Expected risk of misclassification: risk of false negative + risk of false positive

$$E = \int_{A^-} P(+)\rho_+(x)dx + \int_{A^+} P(-)\rho_-(x)dx$$

Which is unknown but can be approximated by

$$E \approx \frac{n_{fn} + n_{fp}}{n}, \quad \text{where } n \text{ is the number of elements in the sample set, } n_{fp} \text{ is the number of false positives in the sample set, } n_{fn} \text{ is the number of false negatives in the sample set.}$$

10.7 Haar Features

Calculating Haar features the naive way:

$$s = \sum_{u=u_0}^{u_0+w-1} \sum_{v=v_0}^{v_0+h-1} g(u, v)$$

The integral image: $I(x, y) := \sum_{u=0}^x \sum_{v_0}^y g(u, v)$ makes calculating s simply a 4-way lookup

on I :

$$s = I(u_0 + w - 1, v_0 + h - 1) - I(u_0 - 1, v_0 + h - 1) + I(u_0 - 1, v_0 - 1) - I(u_0 + w - 1, v_0 - 1)$$

10.8 Ensembles

Sum up all predictions and compare with zero:

$$ensemble(\vec{x}) = \text{sign} \left(\sum_{j=1}^k c_j(\vec{x}) \right)$$

10.9 Boosting

Introduce weight γ_i for each training pattern, applied to e.g. soft margin SVM this yields:

$$\text{minimize } \frac{1}{2} \|\vec{w}\|^2 + C \sum_i \gamma_i \xi_i \quad \text{subject to} \quad d^{(i)}(\langle \vec{x}^{(i)}, \vec{w} \rangle + b) \geq 1 - \xi_i \quad \text{for all } i$$

10.10 AdaBoost

Is an implementation of Boosting. With the training error bounded to

$$\Pi_{t=1}^T \left(2\sqrt{\epsilon_t(1 - \epsilon_t)} \right) \leq \exp \left\{ -2 \sum_{t=1}^T \left(\frac{1}{2} - \epsilon_t \right)^2 \right\}$$

If all $\epsilon_t < \frac{1}{2}$ and $T \rightarrow \infty$ AdaBoost yields a perfect classifier.